



The Team



FSP4

Team - 9




Tejas



Arushi




Raghav






FORECASTING CRYPTOCURRENCY TRENDS

While comparing the effectiveness of
different LSTMs



PROBLEM STATEMENT

Given the volatility of cryptocurrencies, how can we forecast the trend in order to make informed investment decisions?





DATASET

The Alpha Vantage API offers historical and real-time data for stocks, forex, and cryptocurrencies. Several time frames are available ranging from 1-minute bars up to monthly.




Data Collection and Data Cleaning

$$x_{scaled} = \frac{x - x_{min}}{x_{max} - x_{min}}$$




Step 1

Extracted data from Alpha Vantage




Step 2

Organised data as per requirements



Step 3



Feature Scaling using MinMaxScaler



Cryptocurrencies by relevance and market cap:

Enter which cryptocurrency you would like to forecast (Integer between 1 and 10):

1. Bitcoin (BTC)
2. Ethereum (ETH)
3. Binance Coin (BNB)
4. Ripple (XRP)
5. Cardano (ADA)
6. Litecoin
7. Lumen (XLM)
8. EOS.IO (EOS)
9. QTUM (QTUM)
10. TRON (TRX)



```
In [52]: cryptolist_ar = ['BTC', 'ETH', 'BNB', 'XRP', 'ADA', 'LTC', 'XLM', 'EOS', 'QTUM', 'TRX']
```

```
In [53]: import itertools

for (i,j) in zip(cryptolist_ar, crypto_dflist):
    cleaned_data[i + ' close(USD)'] = j['close (USD)']
    cleaned_data[i + ' mktcap(USD)'] = j['market cap (USD)']
```


```
In [54]: cleaned_data.head()
```

```
Out[54]:
```


	timestamp	BTC close(USD)	BTC mktcap(USD)	ETH close(USD)	ETH mktcap(USD)	BNB close(USD)	BNB mktcap(USD)	XRP close(USD)	XRP mktcap(USD)	ADA close(USD)	...	LTC close(USD)	LTC mktcap(USD)
0	2021-04-21	56395.68	1609.784422	2350.71	4.141744e+04	576.7544	426547.695	1.40664	3.694768e+07	1.26773	...	267.23	6.471902
1	2021-04-20	56425.00	72744.482151	2330.03	9.922408e+05	586.3635	5730895.325	1.38501	1.522196e+09	1.26689	...	260.68	2.111225
2	2021-04-19	55633.14	78229.042267	2161.12	8.205923e+05	504.0322	5031325.713	1.30945	1.608074e+09	1.19450	...	261.38	1.973512
3	2021-04-18	56150.01	124882.131824	2235.64	1.475384e+06	481.4367	4468597.460	1.40797	2.048345e+09	1.27693	...	273.36	3.346590
4	2021-04-17	60006.66	58912.256128	2317.60	6.242323e+05	514.6861	2949040.221	1.53896	1.108826e+09	1.36802	...	300.86	2.574520


5 rows x 14 columns

Data Cleaning




EXPLORATORY DATA ANALYSIS





Market Capitalisation of Cryptocurrencies




BIVARIATE STATISTICS

Poor correlation between market cap and closing price.




DATA VISUALIZATION

Pairplot shows a correlation between closing prices of different cryptocurrencies.





Correlation Heat map of Cryptocurrency closing prices




DEEP LEARNING

Long Short-Term Memory Neural
Networks For Time Series
Forecasting







What is an LSTM?



Long short-term memory (LSTM) is an artificial recurrent neural network (RNN) architecture used in the field of deep learning.


- Effective for time series forecasting
- Unlike standard feedforward neural networks, LSTM has feedback connections





Simple LSTM

Single LSTM with hidden dense layer.



Convolutional LSTM



LSTM with convolutional layers.

Stacked LSTMs

Multiple LSTMs working concurrently

Bi-directional LSTM

LSTM which utilizes, both, previous and future data to perform accurate forecasting




```

#Convolutional LSTM
#The shape of input data must be: [samples, timesteps, rows, columns, features]

trainX = trainX.reshape((trainX.shape[0], 1, 1, 1, seq_size))
testX = testX.reshape((testX.shape[0], 1, 1, 1, seq_size))

model = Sequential()
model.add(ConvLSTM2D(filters=64, kernel_size=(1,1), activation='relu', input_shape=(1, 1, 1, seq_size)))
model.add(Flatten())
model.add(Dense(32))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mean_squared_error')
model.summary()

```

```

#Stacked LSTM with 1 hidden dense layer
# The shape of the input data must be: [samples, time steps, features]

#trainX = np.reshape(trainX, (trainX.shape[0], 1, trainX.shape[1]))
#testX = np.reshape(testX, (testX.shape[0], 1, testX.shape[1]))

#model = Sequential()
#model.add(LSTM(50, activation='relu', return_sequences=True, input_shape=(None, seq_size)))
#model.add(LSTM(50, activation='relu'))
#model.add(Dense(32))
#model.add(Dense(1))
#model.compile(optimizer='adam', loss='mean_squared_error')

#model.summary()

```

```

#Bidirectional LSTM
# reshape input to be [samples, time steps, features]
#trainX = np.reshape(trainX, (trainX.shape[0], 1, trainX.shape[1]))
#testX = np.reshape(testX, (testX.shape[0], 1, testX.shape[1]))
#
##For some sequence forecasting problems we may need LSTM to learn
## sequence in both forward and backward directions
#from keras.layers import Bidirectional
#model = Sequential()
#model.add(Bidirectional(LSTM(50, activation='relu'), input_shape=(None, seq_size)))
#model.add(Dense(1))
#model.compile(optimizer='adam', loss='mean_squared_error')
#model.summary()


```

```

#Single LSTM
#trainX = np.reshape(trainX, (trainX.shape[0], 1, trainX.shape[1]))
#testX = np.reshape(testX, (testX.shape[0], 1, testX.shape[1]))

#print('Single LSTM with hidden Dense Layer')
#model = Sequential()
#model.add(LSTM(64, input_shape=(None, seq_size)))
#model.add(Dense(32))
#model.add(Dense(1))
#model.compile(loss='mean_squared_error', optimizer='adam')
#monitor = EarlyStopping(monitor='val_loss', min_delta=1e-3, patience=20,
#                          verbose=1, mode='auto', restore_best_weights=True)
#model.summary()

```





WHY LSTMS?

Reason 1




Allows previous outputs to be used as inputs

Reason 2

Accurate for Time Series forecasting



DATA DRIVEN INSIGHTS



Bi-directional LSTM:

seq_size=10, epochs=50
Train Score: 354.02 RMSE
Test Score: 2421.81 RMSE

seq_size=15, epochs=50
Train Score: 373.15 RMSE
Test Score: 6189.28 RMSE

seq_size=7, epochs=50
Train Score: 358.74 RMSE
Test Score: 2735.18 RMSE

seq_size=10, epochs=100
Train Score: 323.52 RMSE
Test Score: 2654.81 RMSE

Convolutional LSTM:

seq_size=40, epochs=100
Train Score: 324.20 RMSE
Test Score: 3463.53 RMSE

seq_size=10, epochs=50
Train Score: 354.33 RMSE
Test Score: 3981.06 RMSE

Best Case:

seq_size=10, epochs=100
Train Score: 323.22 RMSE
Test Score: 1537.83 RMSE

Stacked LSTMs:

seq_size=10, epochs=100
Train Score: 353.35 RMSE
Test Score: 2899.92 RMSE

seq_size=7, epochs=100
Train Score: 318.81 RMSE
Test Score: 3310.89 RMSE

seq_size=15, epochs=100
Train Score: 392.87 RMSE
Test Score: 4119.80 RMSE

seq_size=10, epochs=50
Train Score: 339.17 RMSE
Test Score: 2191.39 RMSE

seq_size=10, epochs=35
Train Score: 443.97 RMSE
Test Score: 5269.84 RMSE

Single LSTM with Hidden Dense Layer:

seq_size=10, epochs=100
Train Score: 334.80 RMSE
Test Score: 5394.96 RMSE


seq_size=10, epochs=50
Train Score: 347.99 RMSE
Test Score: 3222.51 RMSE

*Comparisons were
performed on data for
Bitcoin (BTC)

Convolutional LSTM Model Summary:

Model: "sequential"

Layer (type)	Output Shape	Param #
conv_lst_m2d (ConvLSTM2D)	(None, 1, 1, 64)	21760
flatten (Flatten)	(None, 64)	0
dense (Dense)	(None, 32)	2080
dense_1 (Dense)	(None, 1)	33
Total params: 23,873		
Trainable params: 23,873		
Non-trainable params: 0		



RESULTS

- Compared RMSE values of various LSTM techniques
- Convolutional LSTM produces reliable results in most cases
- Sequence size between 7-20, #of Epochs 50-100
- Current closing values (uni-variate) are good indicators of future closing values for cryptocurrencies




Beyond The Course



- Neural Networks
- Feature Normalization
- Long Short-Term Memory Neural Networks
- Time series Forecasting
- Overfitting

Team Contributions



Tejas

1. Data Analysis
2. Deep learning
3. Data Cleaning
4. Voice over

Raghav

1. Brainstorming
2. Data cleaning
3. Exploratory Analysis

Arushi

1. Brainstorming
2. Create data frames
3. Slide design & video presentation



References

1. <https://www.3blue1brown.com/>
2. <https://youtu.be/aircAruvnKk>
3. Using a Keras Long Short-Term Memory (LSTM) Model to Predict Stock Prices – kdnuggets
[https://www.kdnuggets.com/2018/11/keras-long-short-term-memory-lstm-model-predict-stock-prices.html#:~:text=\(%2018%3An45%20\)-,Using%20a%20Keras%20Long%20Short%2DTerm%20Memory%20\(LSTM\),Model%20to%20Predict%20Stock%20Prices&text=LSTMs%20are%20very%20powerful%20in,in%20predicting%20its%20future%20price.](https://www.kdnuggets.com/2018/11/keras-long-short-term-memory-lstm-model-predict-stock-prices.html#:~:text=(%2018%3An45%20)-,Using%20a%20Keras%20Long%20Short%2DTerm%20Memory%20(LSTM),Model%20to%20Predict%20Stock%20Prices&text=LSTMs%20are%20very%20powerful%20in,in%20predicting%20its%20future%20price.)
4. <https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-recurrent-neural-networks>